

Write a python program to import and export data using Pandas lib. func's

```

import pandas as pd
airbnb_data = pd.read_csv("listings-austin.csv")
airbnb_data.head()
# Imported Data

# Reading Data from URL
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
col_names = ["sepal-length-in-cm", "sepal-width-in-cm", "petal-length-in-cm", "petal-width-in-cm", "class"]
iris_data = pd.read_csv(url, names=col_names)
iris_data.head()

# Exporting DF to a CSV file
iris_data.to_csv("cleaned-iris-data.csv")

```

Output :

	id	listing_id	scrape_id	cal_host_listings_count	
0					
1					
0	sepal-length-in-cm	width	petal-length-in-cm	width	class
1					

~~21/3/24~~

I Get the Data

→ `fetch_housing_data()`

`import pandas as pd.`

`def load_housing_data(housing_path=HOUSING_PATH):`

`data_path = os.path.join(housing_path,`

`"housing.csv")`

`return pd.read_csv(data_path)`

Quick → `housing = load_housing_data()`

Look `housing.head()`

at the `housing.info()`

data. `housing['Ocean-proximity'].value_counts()`

`housing.describe()`

→ `import matplotlib.pyplot as plt`

`import seaborn as sns`

`housing.hist(bins=50, figsize=(20, 15))`

`plt.show()`

Creating a Test Set

→ `import numpy as np`

`def split_train_test(data, test_ratio=0.2):`

`return data.iloc[train_indices],`

`data.iloc[test_indices]`

`train_set, test_set = split_train_test(data=`

`housing)`
`len(train_set), len(test_set)`

→ housing['income_cat'] = pd.cut(x=housing['median_income'], bins=[0, 1.5, 3, 4.5, 6, np.inf], labels=[1, 2, 3, 4, 5]).

Visualizing

housing['income_cat'].hist()

strat_train_set['income_cat'].value_counts()

len(strat_train_set)

for set in [strat_train_set, strat_test_set]:
 set.drop('income_cat', axis=1, inplace=True)

→ # We tried to generate test set :: it is an imp part of any ML project.

II Discover & Visualize the Data to Gain Insight

→ housing = strat_train_set.copy();

housing.shape

housing.plot(Kind='scatter', x='longitude',
y='latitude')

plt.show()

housing.plot(Kind='scatter', x='longitude',
y='latitude', alpha=0.1)

plt.show()

housing.plot(Kind='scatter', x='longitude',
y='latitude', alpha=.4,
s=housing['population']/100,
label='population', figsize=(10, 7),
c='median_house_value'; cmap=
plt.get_cmap(name='jet'),
colorbar=True)

plt.legend()

housing[['population', 'median_house_value']].

→ Correlations:

corr_matrix = housing.corr()

corr_matrix[['median_house_value']].sort_values(ascending=False)

From pandas.plotting import scatter_matrix
attributes = ['median_house_value', 'median_income',
'total_rooms', 'housing_median_age']

- scatter_matrix (frame = housing[attributes],
figsize = (12, 8))
- plt.show()
- housing.plot (kind = 'scatter', x = 'median_income',
y = 'median_house_value', figsize = (12, 8),
alpha = 0.1)

→ corr_matrix = housing.corr()
corr_matrix[['median_house_value']].sort_values(ascending=False),

III Preparing The Data for ML Algorithms

1. Data Cleaning

II 2. Handling Text and Categorical Attributes

3. Custom transformers

4. Feature Scaling

5. Transformation Pipelines

IV Select and Train a Model.

1. Training & Evaluating on the Training set.
2. Better Evaluation using Cross-Validation

V Fine-Tune Your model.

1. Grid Search
2. Randomized Search
3. Ensemble Methods.
4. Analyze models.
5. Evaluating system

VI Exercises

Ques 1/3/24

Lab Program No. 4Simple Linear Regression

```
df_sal = pd.read_csv('Salary-Data.csv')
df_sal.head()
```

```
df_sal.describe()
```

```
plt.title('Salary Dist plot')
sns.distplot(df_sal['Salary'])
plt.show()
```

```
plt.scatter(df_sal['Years Exp'], df_sal['Sal'],
           color = 'lightcoral')
plt.title('Sal v/s Exp')
plt.xlabel('Years of Exp')
plt.ylabel('Salary')
plt.box(False)
plt.show()
```

```
x = df_sal.iloc[:, :1]
```

```
y = df_sal.iloc[
```

~~X_train, X-test, y-train, y-test =
train-test-split(x, y, test_size = 0.2, random_state = 0)~~

```
regression = LinearRegression()
```

```
regression.fit(x-train, y-train)
```

```
y-pred-test = regression.predict(x-test)
```

```
y-pred-train = regression.predict(x-train)
```

plt.scatter(X_train, y_train, color='lightcoral')

plt.plot(X_train, y_pred_train)

plt.show()

print(f'Coefficient: {regression.coef_}')

print(f'Intercept: {regression.intercept_}')

Multiple Linear Regression

```
df_start = pd.read_csv('50-startups.csv')
df_start.head()
```

```
df_start.describe()
```

```
plt.title('Profile Dist. Plot')
sns.distplot(df_start['Profit'])
plt.show()
```

```
plt.scatter(df_start['R&D Spend'],
            df_start['Profit'], color = 'lightcoral')
plt.show()
```

```
x = df_start.iloc[:, :-1].values
y = df_start.iloc[:, -1].values
```

```
X_train, X_test, y_train, y_test =
train_test_split(X, y, test_size = 0.2,
random_state = 0)
```

~~regression = LinearRegression()~~

~~regression.fit(X_train, y_train)~~

~~y_pred = regression.predict(X-test)~~

```
np.set_printoptions(precision = 2)
```

```
result = np.concatenate((y_pred.reshape(
len(y_pred), 1), y-test.reshape(len(y-test),
1)), 1)
```

~~print(result)~~

~~dfy~~

Building the decision tree (ID3) and apply to classify a new sample

```
from sklearn.datasets import load_iris
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
plot_tree
```

```
import matplotlib.pyplot as plt
```

```
iris = load_iris()
```

```
X = iris.data
```

```
Y = iris.target
```

```
feature_names = iris.feature_names
```

```
target_names = iris.target_names
```

```
X_train, X_test, y_train, y_test =
```

```
train_test_split(X, y, test_size=0.4,
```

```
random_state=42)
```

```
clf = DecisionTreeClassifier(criterion='gini')
```

```
clf.fit(X_train, y_test)
```

```
y_pred = clf.predict(X_test)
```

```
accuracy = clf.score(X_test, y_test)
```

```
print("Accuracy", accuracy).
```

```
plt.figure(figsize=(12, 8))
```

```
plot_tree(clf, feature_names=feature_names)
```

```
class_names=target_names, filled=True)
```

```
plt.show()
```

```
Output: Accuracy = 0.983
```

Date 25/04/24

Lab Program no.5

Build Logistic Regression Model for a given data

```
import pandas as pd.
```

```
from matplotlib import pyplot as plt.  
%matplotlib inline
```

```
df = pd.read_csv("insurance_data.csv")  
df.head()
```

```
plt.scatter(df['age'], df['bought_insurance'],  
marker = '+', color = 'red')
```

```
from sklearn.model_selection import train_test_spli
```

```
X_train, X_test, y_train, y_test =
```

```
train_test_split(df['age'],
```

```
df['bought_insurance'], train_size = 0.8)
```

```
from sklearn.linear_model import LogisticRegression  
model = LogisticRegression()  
model.fit(X_train, y_train)
```

~~y-predicted = model.predict(X-test)~~

~~model.predict_proba(X-test)~~

~~model.score(X-test, y-test)~~

→ Model accuracy : 0.83733332334

Analyzing through Linear regression
model.coef_

→ Coefficient (slope) : [0.1454470877]

model.intercept_

→ Intercept : [-5.447337817]

Analyze through sigmoid function.

import math

def sigmoid(x):

return 1 / (1 + math.exp(-x))

def prediction_function(age):

$z = 0.042 \cdot \text{age} - 153 \# 0.04150133 \sim 0.042 \text{ and } -1.52726963 \sim -1.53$

$y = \text{sigmoid}(z)$

return y

age = 35

prediction_function(age)

→ Prob: 0.48500449838 0.8899

age = 43

prediction_function(age)

→ Prob: 0.568565299077705

Q1 25/4/24

Date 09/05/24

Lab program No. 6

Build KNN classification model for a given dataset.

iris-data = pd.read_csv('Iris.csv')

iris-data = head()

label_mapping = { 'Iris-setosa': 1, 'Iris-versicolor': 2, 'Iris-Virginica': 3 }

iris-data['Species'] = iris-data['Species'].map(label_mapping)

plt.scatter(iris-data['Sepal Length(Cm)'],

iris-data['Sepal Width(Cm)'],

c = iris-data['Species'], cmap = 'viridis')

plt.show()

iris-data['Species'].unique()

array(['Iris-setosa', 'Iris-versicolor',
 'Iris-Virginica'], dtype = object)

iris-data.drop(['Id'], inplace = True, axis = 1)

x = iris-data.drop(['Species'], axis = 1)

y = iris-data['Species']

X-train, X-test, y-train, y-test =
 train-test split(X, y, test_size = 0.4,
 random_state = 42)

Knn = KNeighborsClassifier(n_neighbors = 5),
 Knn.fit(X-train, y-train)

$y_{pred} = \text{Knn.predict}(X_{test})$

accuracy = accuracy_score(y_{test}, y_{pred})
print("Accuracy", accuracy)

Accuracy: 0.983

17/5/24

SVM model

svm-classifier = SVC()

svm-classifier.fit(X_train, y_train)

y-pred = svm-classifier.predict(X-test)

accuracy = accuracy-score(y-test, y-pred)

Accuracy : 0.9333

avg(93.33)

ANN program no. 8

import numpy as np

$X = np.array([[[2, 4], [1, 5], [3, 6]]], dtype=float)$
 $y = np.array([[92], [86], [89]]], dtype=float)$

$X = X / np.array(X, axis=0)$

$y = y / 100$

def sigmoid(x):
 return 1 / (1 + np.exp(-x))

def derivatives_sigmoid(x):
 return x * (1 - x)

for i in range(epoch):

Forward hinp1 = np.dot(x, wh).

Hpropagation. hinp = hinp1 + bh

hlayer_act = sigmoid(hinp).

outinp1 = np.dot(hlayer_act, wout)

outinp = outinp1 + bout

output = sigmoid(outinp)

Backpropagation

EO = y - output

outgrad = derivatives_sigmoid(output)

d_output = EO * outgrad

EHT = ch_output. dot(wout.T)

outgrad_hiddengrad = derivatives_sigmoid(hlayer^{out})

d_hiddengrad = EHT * hiddengrad

updating weights and biases.

$w_{out} += b_{layer_act} \cdot T \cdot \text{dot}(d_output) * l_n$

$w_h += X \cdot T \cdot \text{dot}(d_hidden_layer) * l_n$

print ("Predicted o/p : (n", output)

Predicted o/p : (n, X)

[0.89]

[0.87]

[0.89]. J. (b) Example 1st

((x-1) * x + 1) / 2 = result

(a) binary search job

(x-1) * x = result

: (log2) square as i not

(new x). take gn = log2

11 + 1 = 12 = gn with result

(gn) binary = the result

but > log2 = gn two

(gn) binary = log2

log2 = 3.32

log2 - 1 = 0.3

1.3 * 2 = 2.6 = result

log2 * 0.3 = log2.6

(T binary) take log2.6 - 1 = 1.5

binary without last = binary without last

binary without last = log2.6

Q1 - Random Forest Algo

Random Forest Algo Lab 9 of 9

data = pd.read_csv('Iris.csv')

X = data.drop(columns = ['Species'])

y = data['Species']

X-train, X-test, y-train, y-test =

train-test-split(X, y, test_size = 0.2,
random_state = 42)

rf-classifier = RandomForestClassifier

(n_estimators = 100, random_state = 42).

rf-classifier.fit(X-train, y-train).

y-pred = rf-classifier.predict(X-test).

accuracy = accuracy_score(y-test, y-pred).

print(f'Accuracy: {accuracy * 100:.2f}%')

Accuracy: 100.00%.

Date: 23/5/24

Week 8. Lab program -10

iris =

data = pd.read_csv('Iris.csv')

X = data[['Sepal Length (cm)', 'Sepal Width (cm)',
 'Petal Length (cm)', 'Petal Width (cm)']]

model = KMeans (n_clusters = 3)

model.fit (X)

(S) plt.figure (figsize = (14, 7))

plt.subplot (1, 2, 1)

plt.scatter (X[['Petal Length (cm)'], X[['Petal Width (cm)']]]
 c = model.labels_, cmap = 'viridis',
 s = 100)

plt.show()

NAME: _____

Week 8 Lab program - 11

cancer = load_breast_cancer()

cancer.keys()

df = pd. DataFrame (cancer ['data'],
columns = cancer ['features'])

scaler = StandardScaler()

scaler. fit (df)

scaled_data = scaler. transform (df)

pca = PCA (n_components = 2)

pca. fit (scaled_data)

x_pca = pca. transform (scaled_data)

scaled_data. shape

~~x_pca. shape~~

plt. figure (figsize = (8,6))

plt. scatter (x_pca [:, 0], x_pca [:, 1],

c = cancer ['target'], cmap = 'plasma')

~~30/5/24~~